

## Day 2 - Data Foundations

- Most common (important?) feature of a database is searching (select in SQL)
- the baseline (worst case) for efficiency is linear (loop through all data once)
- **record** - row of a table (collection of values)
- $n \times x$  bytes of memory for records that are  $x$  bytes and  $n$  records
- **Contiguously allocated list**: all stored in a single memory chunk
- **Linked list**: each record needs  $x$  bytes and spaces for 1-12 memory addresses (pointers)
  - memory addresses link records at the end of each record
- Python doesn't have built-in arrays
- arrays - faster for random access, slow for inserting except for at the end
- linked lists - faster for inserting anywhere, slower for random access
- **Binary search** - array of inputs in sorted order, start in middle and split, and do recursively
  - worst case:  $\log(n) + 1$  (base 2)
  - best case: 1 (element is at mid)
- **linear**:
  - worst case is  $n$
  - best case is 1 (first element)
- Generalization: In databases searching for id value is fast, but other columns are slow (because not in sorted order)

## Binary Search Trees

- Use another data structure or an index for columns that need to be queried efficiently
  - binary search tree is how this is done
- Creating/inserting into binary search tree
- Tree traversals
  - pre-order
  - post-order
  - in order
  - level order - go level by level (top to bottom) and left to right within each level
    - temporarily store children of element when processing it (queue)
      - first in first out
      - can use a deque in python (can remove from front or back)
      - this is used to keep a running list of next elements to process
- For homework, question 2:
  - `root = BinaryTreeNode(23)`
  - `root.left = BinaryTreeNode(17)`
- Time complexity Insert is same as linked list, search is as good as sorted binary search
- Duplicate values

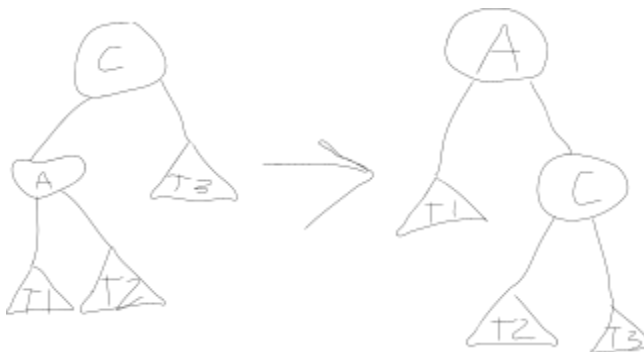
- choose that either always go left, always go right, or are ignored
- Want to minimize height of Tree (make more balanced)
  - need AVL tree for this

## AVL Tree

- approx balanced binary search tree
  - each node has a balance factor
  - $|h(LST) - h(RST)| \leq 1$ 
    - height of left sub tree and right\_sub\_tree should have one or 0 level difference
- Check for imbalance in every node in path to root
- If one node is imbalance, the new node takes spot of imbalance node, and it is rebuilt from there
  - simplification
- 4 cases of imbalance

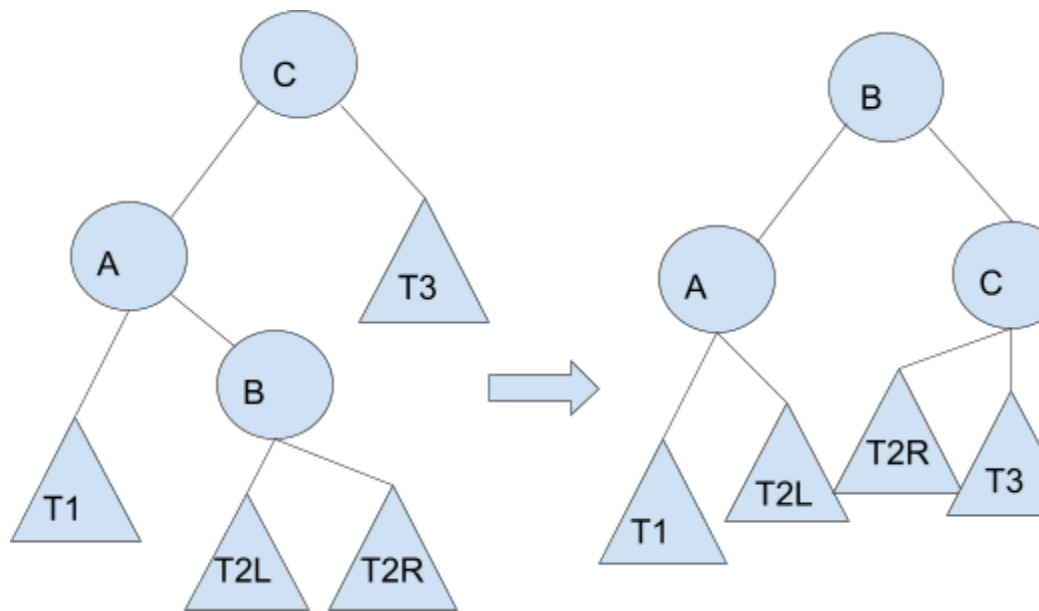
### 1. (LL) left left case

- use a single rotation
- Pseudocode:
  - c.left = a.right
  - a.right = c



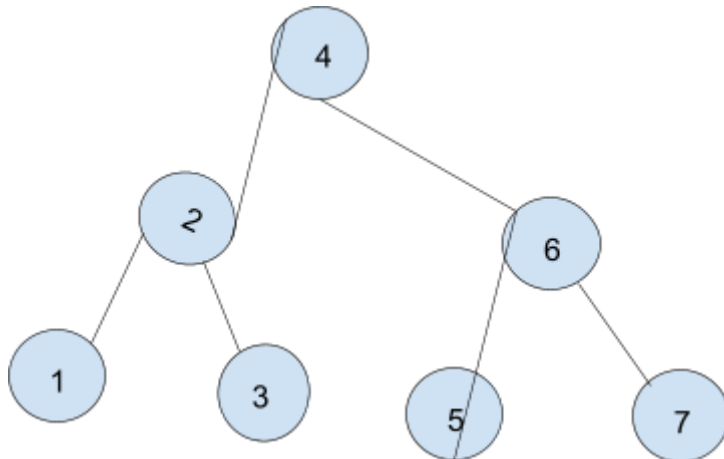
### 2. (LR) left right case

- use a double rotation
- Pseudo code
  - c.left = b.a
  - a.right = b.left
  - b.left = a
  - b.right = c

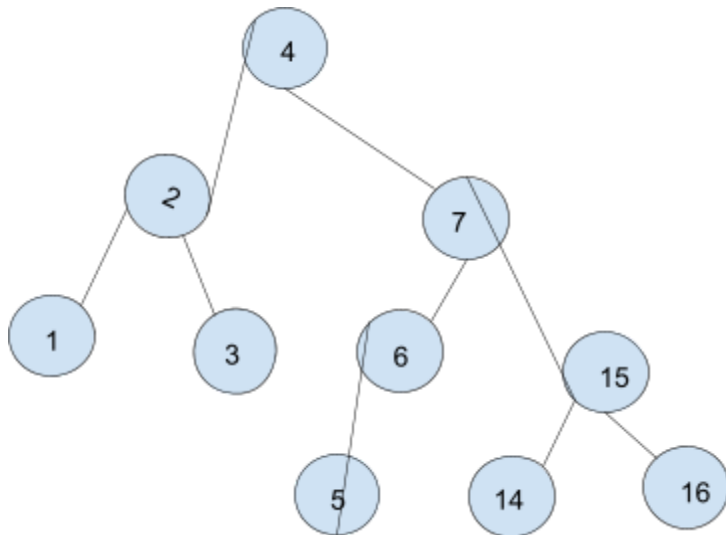


3. (RL) right left case
  - a. mirror of LL (case 1)
  - b. Pseudocode
    - a.right = c.left
    - c.left = a
4. (RR) right right case
  - a. mirror of LR (case 2)
  - b. Pseudocode

- When building out a BST, you start inserting at the root and follow down the tree to the correct position, then check if each position going back up is balanced
  - check the height of left and right subtrees and check the difference  $\leq 1$
  - store height as part of tree node?
    - update height when going back up subtree
  - insert in order [ 3,2,1,4,5,6,7]



Now insert 16, 15, 14



## B+ Trees

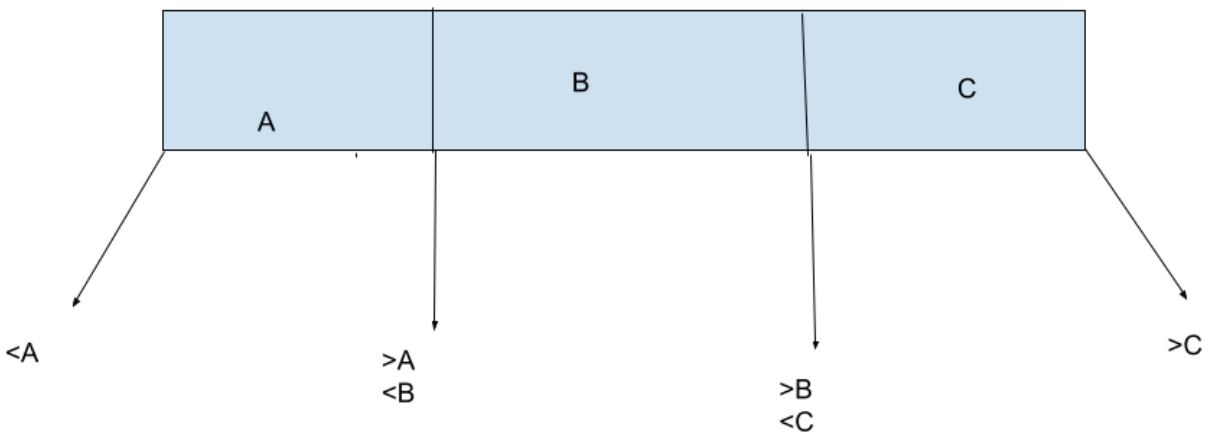
### Day 1

- general computer info
  - Minimize HDD/SDD accesses (slow) Hard drive access
  - CPU reads in chunks from hard drive could be 4kb block
  - things related aren't necessarily next to each other in hard drive
  - AVL trees don't optimize number of blocks read from memory
- each node will have 2 keys and up to 3 children
  - given a node with keys  $k_1$  and  $k_2$ 
    - one children  $< k_1$  and in between  $k_1$  and  $k_2$ , and one  $> k_2$
- data must be sorted
  - this allows for groups of data equal to one block of memory stored together

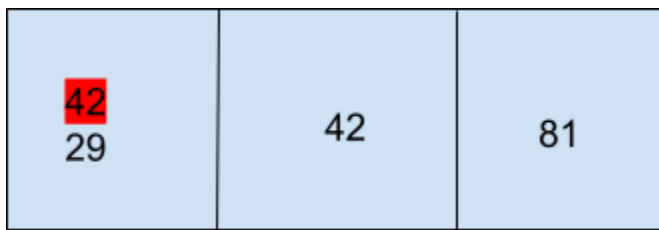
### Day 2

- What is normally used in modern DBs
- goal is to minimize disc accesses
  - optimized for disk-based indexes
- B+ Tree is an m-way tree with order M
  - $M = \text{max \# of keys in each node}$

- $M+1$  = max children of each node
- For  $M=3$



- All nodes (except the root) must be  $\frac{1}{2}$  full min
- root node doesn't have to be  $\frac{1}{2}$  full
  - this only because it starts not half full when you start inserting
- Insertions are always done at leaves
- Leaves are stored as a DLL (doubly linked list)
- Insert 42, 29, 81 in B+ Tree



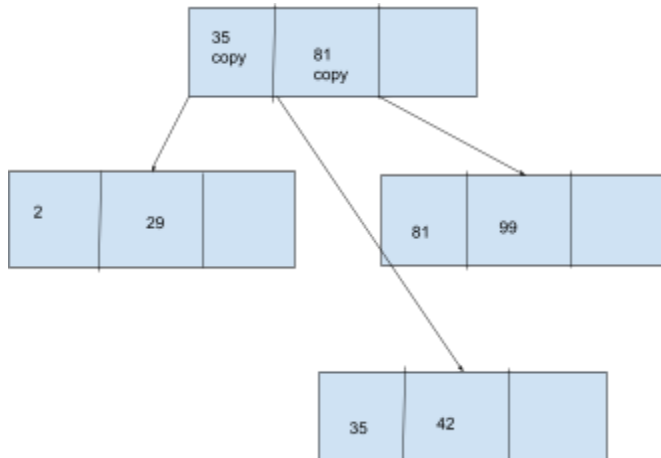
- Start at left then adjust when needed?
- **Internal nodes:** only store keys and pointers to children
- **Leaf nodes:** stores keys and data
- B and B+ trees are different
  - B trees for in-memory indexing
  - B+ trees for disk-based indexing
- Now add 99, 35, 2



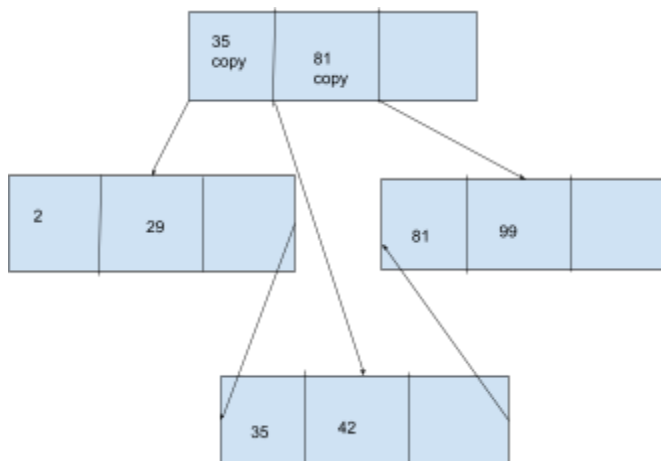
Then:



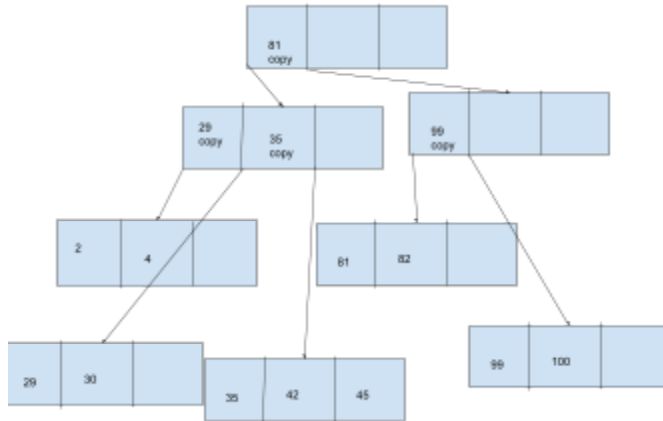
Then



- Always add new node to the right
- leaf nodes are also connected, so final from above is actually



Now add 100, 30, 45, 82, 4



## Hash Tables

- dicts are hash tables, but we can't use them for prac 1
- key values pairs
- every time you give same input then you will get the same output
- # of inserted values = n
- table size = m
- load factor:  $\lambda = \frac{n}{m}$
- constant work for every value k (key) that is being stored
- used a mod operation like  $h(k) = k \bmod i$  where i is an prechosen number (table size)
  - this makes inserting constant time for any k
  - if already an item or multiple at stop, it becomes a list/bucket of in that location
- **(separate chaining)**
- You want to start with a big table size
  - want to keep load factor < .9 then increase size and rehash
- longest chain 5kv pairs, (rehash if greater?)
- can use a premade hashing function for practical (some hashing library?)
- good dispersion (don't want hashes very clustered)
- prebuilt hash function then mod with table size
- 

## Beyond Relational Models

- Benefits of relational models
  - (mostly) standard data model and sql is easy
  - ACID compliance (atomicity, consistency, isolation, durability)

- works well when highly structured
  - easy to work with generally
- RDBMS (relational databases (MS?)) management?
- ways to increase efficiency
  - indexing
  - column vs. row oriented
  - directly controlling storage
  - query optimization
  - caching
  - materialized views
  - precompiled stored procedures
- Transaction processing - a sequence of multiple actions in one unit of work
- Atomicity - transaction is considered on unit (either all fails or all executes)
- Consistency
- Isolation - two transactions happen at same time and will not affect each other
  - Dirty read - a transaction T<sup>1</sup> is able to read a row that has been modified by transaction T<sup>2</sup> that hasn't been committed yet
    - violation of isolation
  - Non-repeatable read
    - two queries execute a single transaction but get different values, because one changed data
  - Often implemented in different ways (can be difficult to ensure without hurting performance)
  - phantom read - transaction adds or deletes rows from the set that another transaction is using
- Durability
  - once a transaction happens, then it is permanent
- why non-relational is sometimes better
  - Not all data needs ACID compliance
  - many data is semi-structured like JSON
  - joins are expensive
  - evolving schemas cause problems with relational databases
  - horizontal scaling, having multiple compute nodes is challenging
- Distributed system
  - a collection of computers working together (if one fails then the others are still ok)
  - no shared clock
- **replications**
  - multiple nodes that have everything
- **sharding**
  - split between different nodes in a cluster
- Major data systems have both sharding and replication
- distributed database can be relational or non-relational
- Network partitioning is inevitable
  - there will be the network failures and system failures



## ● CAP THEOREM

- states that it is impossible for a data store to simultaneously provide more than two out of the following three guarantees
- **Consistency**, every read received the most recent write or error
- **availability** - every request doesn't receive an error, no guarantee of most recent write
- **partition tolerance**- the system can continue to operate despite arbitrary network issues
- CA - is good for relational databases
- CP - is good on nosql databases like mongo, redis
- AP is things like CouchDB, Cassandra, DynamoDB

## NOSQL and KV DBs

- ACID transactions focus of *data safety*
  - considered overly protective in some sense, assumes if something can go wrong it will go wrong
  - use locking data tables to ensure stability/security
  - pessimistic concurrency model
- Optimistic Concurrency - transactions do not obtain locks
  - assumes conflicts are unlikely and assumes things will work out
  - uses timestamp and version number on every table (reads them when changing. check and start and end of each transaction)
  - good for low conflict systems
    - read heavy systems or only updated at night
    - conflicts are solved by rolling back
- Pessimistic is better if a lot of conflicts are expected as roll back and rerunning is not efficient
- Alternative to ACID is BASE
  - Basically Available- data is usually available but sometimes the response will be "failure / unreliable"
    - systems appear to work most of the time
  - Soft State - the state of the system can change over even without updates
    - not all replicas are always the same
  - Eventual Consistency
    - the system will eventually become consistent as long as there is a break in updates
- KEY VALUE databases
  - operate similarly to a hash table
  - very simple data model
  - usually in-memory DB
  - very easily scalable (adding more nodes)

- operate under eventual consistency rather than absolute
- EDA/experiment data
  - can use to store intermediate results from modeling/EDA
- Feature store
  - store frequently accessed data for low latency retrieval
- Model monitoring
  - keep track of metrics on model performance
- shopping cart data
  - tied to the user
- user profiles/preferences
- Storing session info
- Caching layer - most frequent use case for SWE
  - short term storage on top of disk based database

## ● Redis

- remote directory server
- open source, and in memory
- sometimes called a data structure store
- most popular key value database (next is a dynamo DB)
- mostly a Key Value DB but can do other stuff
- developed in 2009 in C++
- very very fast
- only supports lookups by key
- by default provides 16 database numbered 0-15
- commands
  - SET to set something
  - GET to get something
  - SELECT to go to a new DB 0-15
  - INCR increase values by one
  - INCRBY someValue x - increase by x
  - DECR and DECRBY
- HSET - hash set
  - other H-other command for most redis commands
  - HGET, HGETALL, HMGET, HINCRBY
- values in redis can be linked list of string values
  - bidirectional stack/queue?
- Queue like Ops
  - LPUSH to push on to left side of queue
    - Can do RPUSH
  - LPOP and RPOP to pop out
- can have a set as a value
- Set operations
  - SADD - set add
  - SINTER - intersection
  - SDIFF - difference

- SREM
- SRANDMEMBER - random item

## Redis + Python

- redis-py
- import redis
- to connect:
  - `client = redis.Redis(host='localhost', port =6379, db=2, decode_responses=True)`
- similar to redis commands
  - say `r` is the redis option
  - `r.set(key, val)` to set
  - `x = r.get(key)` to get
  - can create pipelines `r.pipeline()`
    - add things to it `pipe.get().set()...`
    - then run `pipe.execute()` to run
    - to decrease call over the network

## Document databases

- type of noSQL database
- usually in a JSON format
- JSON stands for (JavaScript Object Notation)
  - not very space-efficient
  - human-readable
- BSON - Binary JSON
  - what MongoDB uses to store JSON
  - more space efficient, but is still structured like JSON
- XML (eXtensible Markup Language)
  - precursor of JSON
  - sort of like JSON but with HTML-like syntax
    - but can make any tags
  - specific XML query and formatting tools
    - probably don't need to know these but things
- Why doc databases
  - object persistence, in the way that objects are created in OO programming
  - don't need a schema like in a relational DB
  - you can just add new items/keys without doing tons of schema changes
    - eg online shop with different types of items with different attributes

## MongoDB

- all cloud providers have a version of document DBs
  - couchDB, amazon document DB, postgres can sorta do it
- MongoDB - short for humongous database
- owned by google?
  - MongoDB Atlas is their managed DB
- No pre-defined schema (different keys in every object)
- terms
  - collection = table/view
  - document = row
  - field = column
  - embedded document = join
  - reference = foreign key
- has its own query language, that can do similar things
- mongosh → CLI tool
- compass - open source GUI to work with it
- most languages and database apps like datagrip can interact with it
  - PyMongo
- do collection. for commands
  - .find() is like select \*
  - .find({"filter key" : "filter value"}) for where =
- 

## Graph DBs

- can be used to store relations and interactions
  - things like social media
- webs is a graph of pages as nodes and hyperlinks as edges
- set of nodes and vertices
- labels and nodes can both have labels (attributes)
- path: way to get from one node to another node
  - how to get from one to another
- Connected: all nodes are connected
- Weighted: nodes have weights/properties
- **Cyclic**:
- Directed: edges have direction
- Hold **sparse** graph in adjacency list
- hold **dense** graph in adjacency matrix
- **Complete** graphs are when every node is connected
- **Rooted Tree** has root node and no cycles
- **Spanning Tree** get rid of nodes while still being fully connected

- **pathfinding** : shortest path between the nodes
  - algos like djikstras

## Neo4j

- ACID compliant
- language is cypher
- similar to SQL but not the same
- APOC plugin
  - useful graph algos
  - awesome procedures on cypher
- Graph Data Science Plugin
  - similar to APOC
- Docker Compose
  - for managing multiple containers
  -